

# Parallel Quicksort Algorithms in MPI and OpenMP

Marco Tallone

May 13, 2024



DATA SCIENCE &  
ARTIFICIAL INTELLIGENCE



SCIENTIFIC &  
DATA-INTENSIVE COMPUTING

# Introduction

Different parallel versions of the Quicksort algorithm in MPI and OpenMP have been implemented and compared:

- **Task Quicksort** (OpenMP only)
- **Simple Parallel Quicksort**
- **Hyperquicksort**
- **PSRS** (Parallel Sorting by Regular Sampling)

# Serial Quicksort


- 1 Choose a pivot  $\tau$  from the input array  $\mathbf{X}$ ;

7	13	18	2	17	1	14	20	6	10	15	9	3	16	19	4	11	12	5	8
---	----	----	---	----	---	----	----	---	----	----	---	---	----	----	---	----	----	---	---

pivot  $\tau$

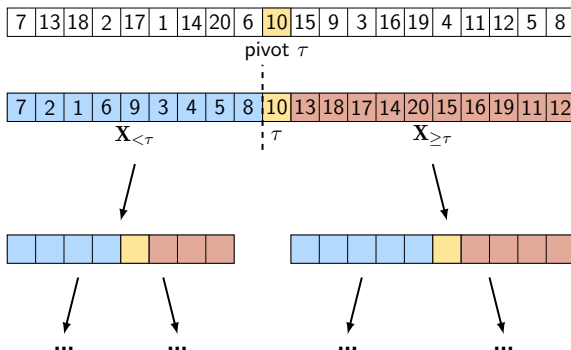
# Serial Quicksort

- 1 Choose a pivot  $\tau$  from the input array  $\mathbf{X}$ ;
- 2 Partition  $\mathbf{X}$  into two subarrays  $\mathbf{X}_{<\tau}$  and  $\mathbf{X}_{\geq\tau}$ ;



# Serial Quicksort

- 1 Choose a pivot  $\tau$  from the input array  $\mathbf{X}$ ;
- 2 Partition  $\mathbf{X}$  into two subarrays  $\mathbf{X}_{<\tau}$  and  $\mathbf{X}_{\geq\tau}$ ;
- 3 Recursively sort  $\mathbf{X}_{<\tau}$  and  $\mathbf{X}_{\geq\tau}$  until  $\text{len}(\mathbf{X}) \leq 2$ .



# Serial Quicksort

## Serial Quicksort

```
1  function serial_qsort(X):
2      if len(X) > 2:
3           $\tau \leftarrow \text{choose\_pivot}(X)$ 
4           $X_{<\tau}, X_{\geq\tau} \leftarrow \text{partition}(X, \tau)$ 
5
6
7          serial_qsort( $X_{<\tau}$ )
8
9
10         serial_qsort( $X_{\geq\tau}$ )
11     else:
12         if len(X) == 2 and  $X[0] > X[1]$ :
13             swap( $X[0], X[1]$ )
```

# Task Quicksort

## OpenMP Task Quicksort

```
1  function omp_task_qsort(X):
2      if len(X) > 2:
3           $\tau \leftarrow \text{choose\_pivot}(X)$ 
4           $X_{<\tau}, X_{\geq\tau} \leftarrow \text{partition}(X, \tau)$ 
5
6          #pragma omp task
7          omp_task_qsort( $X_{<\tau}$ )
8
9          #pragma omp task
10         omp_task_qsort( $X_{\geq\tau}$ )
11     else:
12         if len(X) == 2 and  $X[0] > X[1]$ :
13             swap( $X[0], X[1]$ )
```



This function needs to be called **inside a** `#pragma omp parallel` **parallel region** to work properly.

# Simple Parallel Quicksort: MPI implementation


- 1 Split  $X$  in  $P$  chunks;

$P_0$					$P_1$					$P_2$					$P_3$				
7	13	18	2	17	1	14	20	6	10	15	9	3	16	19	4	11	12	5	8




## Simple Parallel Quicksort: MPI implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Choose **one** pivot  $\tau$  and broadcast it;




# Simple Parallel Quicksort: MPI implementation

- 1 Split  $\mathbf{X}$  in  $P$  chunks;
- 2 Choose **one** pivot  $\tau$  and broadcast it;
- 3 Local partitioning in  $\mathbf{X}_{<\tau}^i$  and  $\mathbf{X}_{\geq\tau}^i$ ;




# Simple Parallel Quicksort: MPI implementation

- 1 Split  $\mathbf{X}$  in  $P$  chunks;
- 2 Choose **one** pivot  $\tau$  and broadcast it;
- 3 Local partitioning in  $\mathbf{X}_{<\tau}^i$  and  $\mathbf{X}_{\geq\tau}^i$ ;
- 4 “Low” partitions  
 $\rightarrow \text{rank} < P/2$ ,  
 “High” partitions  
 $\rightarrow \text{rank} \geq P/2$ ;



# Simple Parallel Quicksort: MPI implementation

- 1 Split  $\mathbf{X}$  in  $P$  chunks;
- 2 Choose **one** pivot  $\tau$  and broadcast it;
- 3 Local partitioning in  $\mathbf{X}_{<\tau}^i$  and  $\mathbf{X}_{\geq\tau}^i$ ;
- 4 “Low” partitions  
 $\rightarrow \text{rank} < P/2$ ,  
 “High” partitions  
 $\rightarrow \text{rank} \geq P/2$ ;
- 5 Recursively sort the chunks.




# Simple Parallel Quicksort: OpenMP implementation

1 Split  $X$  in  $P$  chunks;

$P_0$				$P_1$				$P_2$				$P_3$				$P_4$			
7	13	18	2	17	1	14	20	6	10	15	9	3	16	19	4	11	12	5	8

# Simple Parallel Quicksort: OpenMP implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Choose **one** pivot  $\tau$  (**shared** variable);




# Simple Parallel Quicksort: OpenMP implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Choose **one** pivot  $\tau$   
(**shared** variable);
- 3 Local partitioning in  $X_{<\tau}^i$  and  $X_{\geq\tau}^i$ ;

$P_0$				$P_1$				$P_2$				$P_3$				$P_4$			
7	13	18	2	17	1	14	20	6	10	15	9	3	16	19	4	11	12	5	8
pivot $\tau$																			
2	7	18	13	1	17	14	20	6	10	15	9	3	4	19	16	5	12	11	8

# Simple Parallel Quicksort: OpenMP implementation


- 1 Split  $X$  in  $P$  chunks;
- 2 Choose **one** pivot  $\tau$  (**shared** variable);
- 3 Local partitioning in  $X_{<\tau}^i$  and  $X_{\geq\tau}^i$ ;
- 4  $L \leftarrow \#$  elements in low partitions,  
 $H \leftarrow \#$  elements in high partitions;





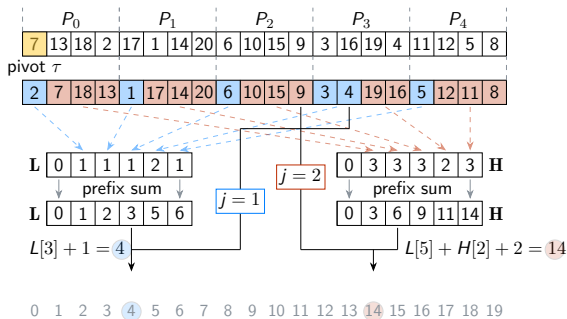
# Simple Parallel Quicksort: OpenMP implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Choose **one** pivot  $\tau$  (**shared** variable);
- 3 Local partitioning in  $X_{<\tau}^i$  and  $X_{\geq\tau}^i$ ;
- 4  $L \leftarrow \#$  elements in low partitions,  
 $H \leftarrow \#$  elements in high partitions;
- 5  $L$  and  $H$  prefix sums;



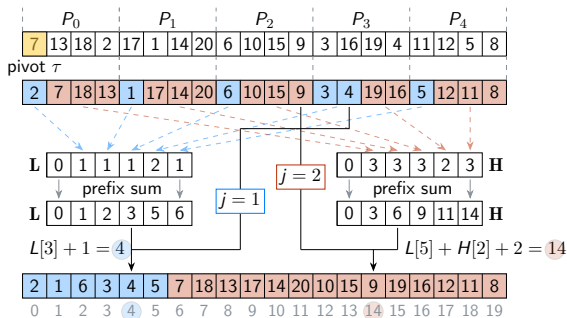
# Simple Parallel Quicksort: OpenMP implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Choose **one** pivot  $\tau$  (**shared** variable);
- 3 Local partitioning in  $X_{<\tau}^i$  and  $X_{\geq\tau}^i$ ;
- 4  $L \leftarrow \#$  elements in low partitions,  
 $H \leftarrow \#$  elements in high partitions;
- 5  $L$  and  $H$  prefix sums;
- 6 Update a **shared** `indexes` array;



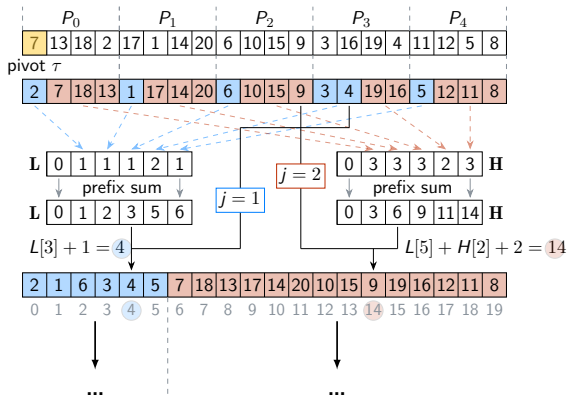
# Simple Parallel Quicksort: OpenMP implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Choose **one** pivot  $\tau$  (**shared** variable);
- 3 Local partitioning in  $X_{<\tau}^i$  and  $X_{\geq\tau}^i$ ;
- 4  $L \leftarrow \#$  elements in low partitions,  
 $H \leftarrow \#$  elements in high partitions;
- 5  $L$  and  $H$  prefix sums;
- 6 Update a **shared** `indexes` array;
- 7 **Serially reorder** according to `indexes`;



# Simple Parallel Quicksort: OpenMP implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Choose **one** pivot  $\tau$  (**shared** variable);
- 3 Local partitioning in  $X_{<\tau}^i$  and  $X_{\geq\tau}^i$ ;
- 4  $L \leftarrow \#$  elements in low partitions,  
 $H \leftarrow \#$  elements in high partitions;
- 5  $L$  and  $H$  prefix sums;
- 6 Update a **shared** `indexes` array;
- 7 **Serially reorder** according to `indexes`;
- 8 Repeat recursively.




# Hyperquicksort

- 1 Split  $\mathbf{X}$  in  $P$  chunks;
- 2 **Each process locally sorts its chunk**
- 3 **One process chooses the median of its chunk as pivot  $\tau$**   
and broadcasts it;
- 4 Local partitioning in  $\mathbf{X}_{<\tau}^i$  and  $\mathbf{X}_{\geq\tau}^i$ ;
- 5 “Low” partitions  $\rightarrow \text{rank} < P/2$ ,  
“High” partitions  $\rightarrow \text{rank} \geq P/2$ ;
- 6 Recursively sort the chunks.

# PSRS: MPI Implementation

- 1 Split  $X$  in  $P$  chunks;



# PSRS: MPI Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;

$P_0$								$P_1$								$P_2$							
7	13	18	2	17	1	14	20	6	10	15	9	3	16	19	4	11	12	5	8				
1	2	7	13	14	17	18	3	6	9	10	15	16	20	4	5	8	11	12	19				

# PSRS: MPI Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;
- 3 Each selects  $P$  samples:

$$\frac{i \cdot n}{P^2}, \quad i = 0, 1, \dots, P - 1$$

$P_0$										$P_1$										$P_2$									
7	13	18	2	17	1	14	20	6	10	15	9	3	16	19	4	11	12	5	8										
1	2	7	13	14	17	18	3	6	9	10	15	16	20	4	5	8	11	12	19										




# PSRS: MPI Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;
- 3 Each selects  $P$  samples:

$$\frac{i \cdot n}{P^2}, \quad i = 0, 1, \dots, P - 1$$

- 4 One process:
  - Gathers and sorts all samples;
  - Regularly picks  $P - 1$  pivots;
  - Broadcasts pivots.




# PSRS: MPI Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;
- 3 Each selects  $P$  samples:

$$\frac{i \cdot n}{P^2}, \quad i = 0, 1, \dots, P - 1$$

- 4 One process:
  - Gathers and sorts all samples;
  - Regularly picks  $P - 1$  pivots;
  - Broadcasts pivots.
- 5 Local partition according to pivots;




# PSRS: MPI Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;
- 3 Each selects  $P$  samples:

$$\frac{i \cdot n}{P^2}, \quad i = 0, 1, \dots, P - 1$$

- 4 One process:
  - Gathers and sorts all samples;
  - Regularly picks  $P - 1$  pivots;
  - Broadcasts pivots.
- 5 Local partition according to pivots;
- 6 **All-to-all** exchange:
  - $P_i$  sends  $j^{\text{th}}$  partition to  $P_j$ ;




# PSRS: MPI Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;
- 3 Each selects  $P$  samples:

$$\frac{i \cdot n}{P^2}, \quad i = 0, 1, \dots, P - 1$$

- 4 One process:
  - Gathers and sorts all samples;
  - Regularly picks  $P - 1$  pivots;
  - Broadcasts pivots.
- 5 Local partition according to pivots;
- 6 **All-to-all** exchange:  
 $P_i$  sends  $j^{\text{th}}$  partition to  $P_j$ ;
- 7 Local sorting.



# PSRS: OpenMP Implementation

- 1 Split  $X$  in  $P$  chunks;

$P_0$							$P_1$							$P_2$					
7	13	18	2	17	1	14	20	6	10	15	9	3	16	19	4	11	12	5	8

# PSRS: OpenMP Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;

$P_0$							$P_1$							$P_2$						
7	13	18	2	17	1	14	20	6	10	15	9	3	16	19	4	11	12	5	8	
1	2	7	13	14	17	18	3	6	9	10	15	16	20	4	5	8	11	12	19	

# PSRS: OpenMP Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;
- 3 Each selects  $P$  samples:  
→ **shared** variable;

$P_0$							$P_1$							$P_2$						
7	13	18	2	17	1	14	20	6	10	15	9	3	16	19	4	11	12	5	8	
1	2	7	13	14	17	18	3	6	9	10	15	16	20	4	5	8	11	12	19	

# PSRS: OpenMP Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;
- 3 Each selects  $P$  samples:  
→ **shared** variable;
- 4 Master thread:
  - Sorts all samples;
  - Regularly picks  $P - 1$  pivots;
  - Writes pivots in shared array.


$P_0$							$P_1$							$P_2$						
7	13	18	2	17	1	14	20	6	10	15	9	3	16	19	4	11	12	5	8	
1	2	7	13	14	17	18	3	6	9	10	15	16	20	4	5	8	11	12	19	

Regularly sampled pivots:  $\tau_0 = 7$ ;  $\tau_1 = 12$




# PSRS: OpenMP Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;
- 3 Each selects  $P$  samples:  
→ **shared** variable;
- 4 Master thread:
  - Sorts all samples;
  - Regularly picks  $P - 1$  pivots;
  - Writes pivots in shared array.
- 5 Local partition according to pivots;



# PSRS: OpenMP Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;
- 3 Each selects  $P$  samples:  
→ **shared** variable;
- 4 Master thread:
  - Sorts all samples;
  - Regularly picks  $P - 1$  pivots;
  - Writes pivots in shared array.
- 5 Local partition according to pivots;
- 6  $(P + 1) \times (P + 1)$  matrix  $M$ ;




0	0	0	0
0	2	2	2
0	1	2	2
0	4	3	2

$M$

# PSRS: OpenMP Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;
- 3 Each selects  $P$  samples:  
→ **shared** variable;
- 4 Master thread:
  - Sorts all samples;
  - Regularly picks  $P - 1$  pivots;
  - Writes pivots in shared array.
- 5 Local partition according to pivots;
- 6  $(P + 1) \times (P + 1)$  matrix  $M$ ;



0	0	0	0
0	2	4	6
0	1	3	5
0	7	4	9



M

→

prefix sum

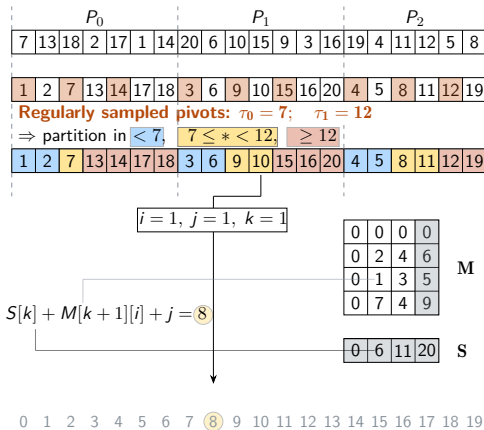
# PSRS: OpenMP Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;
- 3 Each selects  $P$  samples:  
→ **shared** variable;
- 4 Master thread:
  - Sorts all samples;
  - Regularly picks  $P - 1$  pivots;
  - Writes pivots in shared array.
- 5 Local partition according to pivots;
- 6  $(P + 1) \times (P + 1)$  matrix  $M$ ;
- 7 Prefix sum  $S$  of last column;



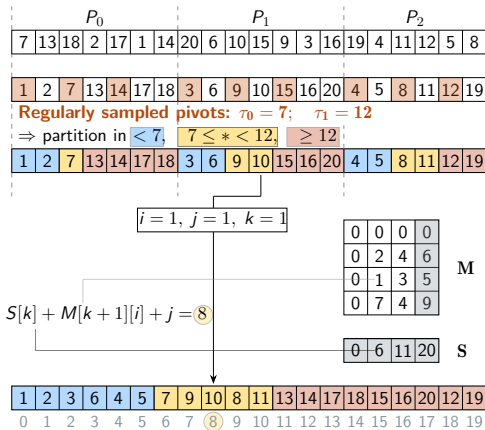
# PSRS: OpenMP Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;
- 3 Each selects  $P$  samples:  
→ **shared** variable;
- 4 Master thread:
  - Sorts all samples;
  - Regularly picks  $P - 1$  pivots;
  - Writes pivots in shared array.
- 5 Local partition according to pivots;
- 6  $(P + 1) \times (P + 1)$  matrix  $M$ ;
- 7 Prefix sum  $S$  of last column;
- 8 Update the `indexes` array;



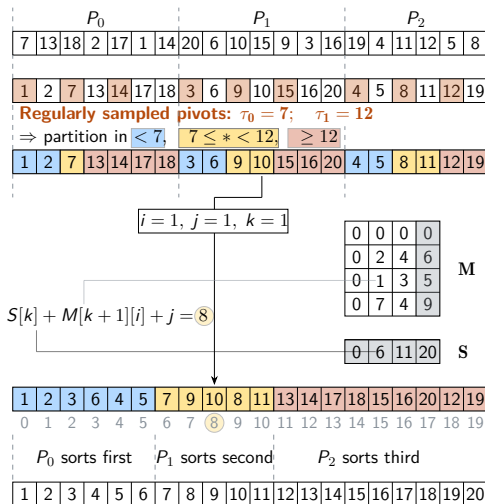
# PSRS: OpenMP Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;
- 3 Each selects  $P$  samples:  
→ **shared** variable;
- 4 Master thread:
  - Sorts all samples;
  - Regularly picks  $P - 1$  pivots;
  - Writes pivots in shared array.
- 5 Local partition according to pivots;
- 6  $(P + 1) \times (P + 1)$  matrix  $M$ ;
- 7 Prefix sum  $S$  of last column;
- 8 Update the `indexes` array;
- 9 **Serially reorder** elements according to `indexes`;



# PSRS: OpenMP Implementation

- 1 Split  $X$  in  $P$  chunks;
- 2 Local chunk **sorting**;
- 3 Each selects  $P$  samples:  
→ **shared** variable;
- 4 Master thread:
  - Sorts all samples;
  - Regularly picks  $P - 1$  pivots;
  - Writes pivots in shared array.
- 5 Local partition according to pivots;
- 6  $(P + 1) \times (P + 1)$  matrix  $M$ ;
- 7 Prefix sum  $S$  of last column;
- 8 Update the `indexes` array;
- 9 **Serially reorder** elements according to `indexes`;
- 10 Each thread sorts from  $S[i]$  to  $M[i + 1][P]$



# Scaling Analysis





## Strong Scaling

Measuring how the execution time  $t$  varies with the number of processors  $P$  for a fixed total workload.





## Weak Scaling

Measuring how the execution time  $t$  varies with the number of processors  $P$  for a fixed workload per processor.


### MPI Implementation


-  Strong Scaling  $t \sim P$ ;
-  Strong Scaling  $S_p \sim P$ ;
-  Weak Scaling  $t \sim P$ ;
-  Weak Scaling  $\varepsilon \sim P$ .


### OpenMP Implementation

-  Strong Scaling  $t \sim P$ ;
-  Strong Scaling  $S_p \sim P$ ;
-  Weak Scaling  $t \sim P$ ;
-  Weak Scaling  $\varepsilon \sim P$ .




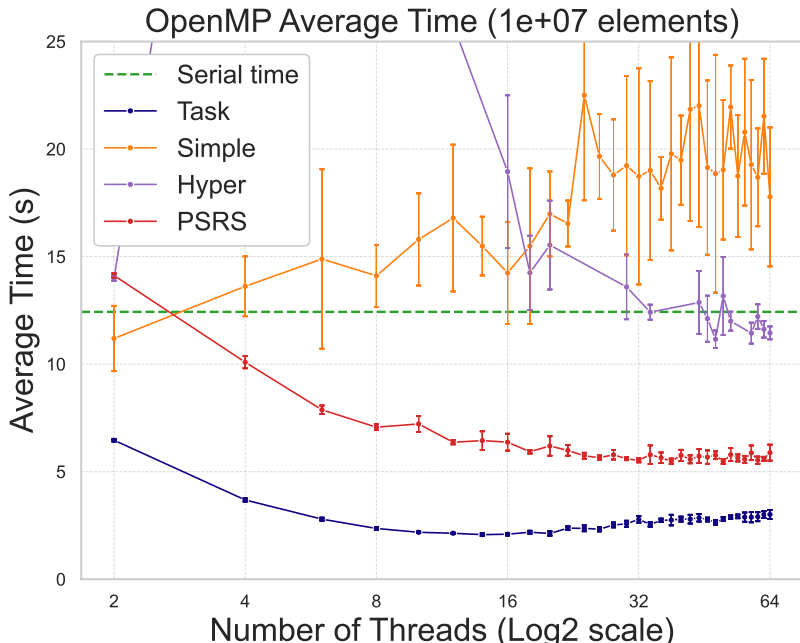







## MPI Efficiency (1e+06 elements per process)







## OpenMP Speedup (1e+07 elements)



## OpenMP Average Time (1e+06 elements per thread)



## OpenMP Efficiency (1e+06 elements per thread)





## References



[AMD.](#)

Epyc 7h12 specifications, 2023.



[OpenMP Architecture Review Board.](#)

Openmp application program interface, 2023.



[MPI Forum.](#)

Mpi: A message-passing interface standard, 2023.



[Ananth Grama, Anshul Gupta, George Karypis, and Vipin Kumar.](#)

*Introduction to Parallel Computing.*

Addison-Wesley, 2nd edition, 2003.



[Frank Nielsen.](#)

*Introduction to HPC with MPI for Data Science.*

Springer, 2016.



[AREA Science Park.](#)

Orfeo documentation, 2023.